

Phase 2 — Agent Core

A summary of the work completed

27 July 2026

What this phase covered

With the database in place, this phase focused on making sure the three outreach accounts can each run safely, independently, and without ever interfering with one another. This is the foundation that protects the accounts themselves before any real activity begins on them.

Hussein built a system that keeps track of all three accounts, knows the daily schedule each one should run on, and decides automatically which accounts are due to run at any given moment — skipping any account that is currently paused, and never running the same account twice in one day.

Keeping every account completely separate

The most important part of this phase was building true isolation between the three accounts. Each account now runs inside its own private, self-contained browser session — nothing it does, no cookie it stores, no login it holds, can ever cross over into another account's session. This was verified directly: three accounts were run at the same time, each was given its own unique marker, and every single one came back seeing only its own data, with zero overlap.

Each account's session is also remembered between runs, the same way a person's browser stays logged in from one day to the next, rather than starting fresh and having to log in again every single time — which is exactly the kind of behaviour that draws unwanted attention on real platforms.

Watching for trouble, without overreacting

A monitoring system was built to watch every account's session for signs of a problem — anything from the page failing to load at all, to being blocked, to being shown a login challenge or a verification screen. When something goes wrong, the specific type of problem and the exact reason are recorded, and that one account is safely paused. Nothing else happens automatically: no account is ever brought back online, and no workload is ever shifted onto the other accounts, without Hussein deciding that himself. This was tested

directly by deliberately causing a failure and confirming the system caught it, recorded the precise reason, and paused only the affected account — while leaving the decision of what to do next entirely in Hussein's hands.

Running it for real

Everything above was tested against the live system, not just in theory. A full run was triggered for all three accounts, and each one completed successfully, with an accurate record of exactly when it started and when it finished. A second test deliberately introduced a failure, and the system correctly detected it, explained why, and paused the right account — then Hussein reset everything back to a clean, ready state afterward.

Result: every check passed. All three accounts can now run safely, in complete isolation from one another, with automatic problem detection and no risk of one account's activity ever affecting another.

Phase 2 of the Nexaris AI Outreach Agent — completed and verified, 27 July 2026.