

AI Outreach Agent

Complete Final Specification

Every decision, concept, task, dashboard detail, setting, and rule

from A to Z — verified against the full planning conversation

Nexaris — Hussein

This is the final, fully verified specification document for the Nexaris AI outreach agent. It was produced after a complete review of the entire planning conversation from the first message to the last, cross-checked against two earlier draft PDFs to ensure nothing was missed. Every agent task, every dashboard screen, every metric, every editable setting, every concept discussed, every decision made, and every rule agreed is captured here.

1. Original Requirements — What Was Asked For

These are the exact requirements from the very first message of the planning conversation. Everything that follows in this document traces back to these original points.

What was asked for, verbatim in concept

- An AI agent for a marketing agency — accessible on web/app or WhatsApp
- Goes to Instagram and searches for 20 clients per day, related to e-commerce (later expanded to any business type)
- Analyzes their profile and all details related to their business
- Generates a suitable message to send as a DM
- Rates if it is a good deal or not AND WHY (became: lead scoring 1-10 with written reasoning)
- Sends feedback to us (became: WhatsApp alerts and daily summary)
- A section showing previous data: how many clients we reached, how many deals, how many meetings, how many hot clients, warm or cold (became: the analytics/KPI dashboard + Client History section + manual deal and meeting logging in pipeline)

How scope evolved during planning

- Platform expanded from Instagram-only to LinkedIn + Instagram + WhatsApp
- Target expanded from e-commerce only to any business type
- CRM, pipeline, follow-up, and re-engagement were added as core requirements
- Advanced features (meetings, proposals, call summaries, customer support, campaigns) were explored in detail then deliberately cut — outreach agent with CRM only is the current scope
- Agent brain confirmed as Claude API (Haiku for analysis, Sonnet for messages)
- Database confirmed as Supabase free tier (2 users only)
- Platform confirmed as responsive PWA — works on laptop and Mahmoud's iPhone

2. Platforms & Channel Strategy

Three channels with three different rules

- LinkedIn — fully automated. Agent finds leads AND sends messages automatically.
- WhatsApp — fully automated but conditional: triggers automatically whenever the agent finds a public WhatsApp number on ANY lead's profile (Instagram bio, LinkedIn profile, or website). Applies regardless of which platform the lead was discovered on. This is an additional channel on top of LinkedIn, not a replacement.
- Instagram — find and analyze ONLY. Agent discovers leads, scores them, and generates their message, but does NOT send. Instagram's official API does not permit automated cold outreach and automating via a browser session violates their ToS and risks bans. Messages are placed in a dedicated dashboard section for a human to send manually.

3. Target Leads & Niche

Who the agent targets

- NOT restricted to e-commerce — any real business that could benefit from marketing or AI services: restaurants, gyms, clinics, law firms, real estate agencies, fashion brands, tech companies, etc.
- Fully configurable per campaign — niche/industry/business type is editable from the dashboard at any time without touching code.
- Country and market based — Hussein chooses the exact location and target market from the dashboard. Fully geography-configurable, not hardcoded to Lebanon or any market.

4. Accounts & Daily Volume

Account pool

- 3 accounts managed by ONE agent (not 3 separate agents — one agent with an account pool)
- All 3 accounts were originally created from the same phone/laptop — creates a weak pre-existing link. Going forward each account gets its own dedicated proxy IP and the personal phone must not log into these accounts again. See Section 11 for full proxy plan.
- Daily volume per account: 20 Instagram profiles + 30 LinkedIn profiles = 50/account
- Total daily capacity: 60 Instagram + 90 LinkedIn = 150 leads/day across all 3 accounts
- These are STEADY-STATE targets — not launch-day numbers. See warm-up plan in Section 11.

5. Scheduling & Execution

Full automatic daily run — no manual launch needed

- Agent runs fully automatically every day.
- Base start time is 9 AM but each of the 3 accounts has its OWN individual run time, staggered through the day — editable per account from the dashboard.
- The agent tracks and displays the exact FINISH TIME of each daily run in the dashboard so the team always knows how long the cycle took.
- At the start of every run: loads all dashboard settings + the full list of previously contacted leads so no one is ever contacted twice unintentionally.

6. Agent Tasks — Complete Workflow (Every Task in Order)

Every single task the agent performs each daily run, in execution order.

Task 1 — Boot and load configuration

- Starts at the scheduled time for each account.
- Loads all dashboard settings: niche, location, language, daily targets, message style, style duration, proxy credentials per account.
- Loads the full previously-contacted leads list (deduplication — no one contacted twice).
- Loads the re-contact list: leads eligible for a second touch based on the gap Hussein set.

Task 2 — Account health check and session open

- Selects the account whose scheduled time has arrived.
- Checks the account's session health.
- If WARNING detected: records the warning type AND reason (specific, not generic — e.g. 'Login challenge detected', 'CAPTCHA triggered', 'Daily limit reached', 'Session expired'). Shows in dashboard. Pauses account. Waits for Hussein to decide whether to redistribute quota — NOT automatic.
- If healthy: opens a dedicated browser session routed exclusively through that account's assigned proxy IP. Each account's session is fully isolated — IPs never mix between accounts.

Task 3 — LinkedIn discovery

- Searches LinkedIn using configured niche, industry, company size, location, and activity filters.
- Targets 30 profiles per account per day (editable from dashboard).
- Reads: company name, description, headcount, website, recent posts, public contact info.
- Prioritizes recently active accounts over dormant ones.
- Agent reasons about search results — if a search returns irrelevant results, it adapts its search terms rather than blindly collecting bad leads (agentic behavior, not fixed automation).

Task 4 — Instagram discovery

- Searches Instagram using hashtags, explore pages, and location tags matched to niche.
- Targets 20 profiles per account per day (editable from dashboard).
- Reads: bio, follower count, post count, engagement on recent posts, story signals.
- Does NOT send any messages at this stage — find and read only.
- If the agent detects an account seems inactive or irrelevant, it skips and moves on.

Task 5 — Business qualification

- For every profile found on either platform: is this a real business that could use marketing or AI services?
- SKIP: personal accounts, inactive pages, no clear product/service, obviously wrong niche.
- PASS: qualified leads move forward immediately.
- Agent makes a reasoning decision here — not a keyword filter.

Task 6 — Deep profile analysis

- Collects the full company dossier for every qualified lead:
- Company size (estimated headcount)
- Estimated revenue tier (from audience size, ads signals, website quality)
- Industry (from content, bio, website)
- Website analysis: design quality, booking flow, CTA, load speed
- Advertising activity: whether they run paid ads (public ad library, visible story ads)
- All social media platforms they are active on
- Full weak points list: low engagement, no website, manual booking, slow response, inconsistent branding, no follow-up system, poor content quality
- AI opportunities: specific areas where automation would help their business
- ALL of this is stored and shown in full in the dashboard — not summarized or hidden.
- The 'surface-level signals only' idea was explicitly REMOVED during planning.

Task 7 — Founder detection

- Actively searches each profile bio for founder signals using phrases including (but not limited to): 'founder of [business name]', 'co-founder', 'CEO & founder', 'established by', 'created by', 'owner of', or any wording that identifies a named person as the founder or owner of the business.
- If a founder is identified: the dashboard shows TWO flags on the same lead card — (1) a message was sent to the company profile, and (2) a founder was found with their name and any available contact signals.
- A WhatsApp alert is sent to both recipient numbers when a founder is found alongside a company message.
- Manual contact with the founder is NOT automated — the team decides how and whether to reach out to the founder separately.

Task 8 — WhatsApp number detection

- On every qualified lead (Instagram AND LinkedIn), the agent checks for a public WhatsApp number in: the bio, profile description, website link, or contact section.
- If found: the WhatsApp outreach workflow fires automatically for that lead, sending a message via WhatsApp IN ADDITION TO (not instead of) the primary channel message.
- Only publicly listed business numbers are used — never scraped or guessed.

Task 9 — Lead scoring and rating

- Every qualified lead is scored 1 to 10 — directly answering the original requirement to 'rate if it is a good deal or not AND WHY'.
- Score components: fit (how well their business matches Nexaris services), pain (number and severity of solvable problems found), budget potential (revenue tier signals).
- Label assigned: Hot (8-10), Warm (5-7), Cold (1-4).
- A written reasoning line explains the score — the WHY is always present.
- Score, label, and reasoning are stored permanently in Client History.

Task 10 — Save to Client History (permanent record)

- Before any message is generated, the full lead record is saved to Client History.
- Client History is a permanent, dedicated section — every lead ever analyzed lives here with full info, regardless of whether outreach was ever sent.
- This directly answers the original requirement: 'a section showing previous data — how many clients did we reach, how many hot/warm/cold.'

Task 11 — Message generation (AI-written, not human-written)

- CONFIRMED: the Claude API (Sonnet model) generates each message — Hussein does NOT write the messages. The AI writes them, personalized per lead.
- Messages must sound fully human in tone — not obviously AI-generated.
- Personalized greeting: uses whichever name is available — business name, team name, or founder name. Example: 'Hi [name]...'
- Message content references the specific weak points and/or opportunities found for that specific business.
- Personalization line example (from the planning conversation): 'I noticed you are spending on Meta ads, but your follow-up is manual and your booking process takes over 12 hours.' — this is the STYLE of personalization the agent should produce.
- MESSAGE STYLE (chosen in dashboard) controls the approach:
 - Direct style: explicitly names the weak points found
 - Discovery style: softer and curiosity-driven, does not state weak points outright
 - Style stays consistent for the configured STYLE DURATION (e.g. one week), then rotates.
- Both style choice and style duration are editable from the dashboard at any time.
- Model split: Haiku for analysis/scoring (cost-efficient), Sonnet for message generation (better writing quality, higher reply rates).

Task 12 — Queue for Mahmoud's approval

- All generated messages for the day are placed in a dedicated approval queue in the dashboard before anything is sent.
- Mahmoud reviews each message, can edit it, and approves or holds it.
- NOTHING sends — on LinkedIn, WhatsApp, or for Instagram manual-send — until Mahmoud gives daily approval.
- This is a permanent hard gate. No bypass toggle exists.

Task 13 — Route and send after approval

- LinkedIn: agent sends the approved message automatically.
- WhatsApp: agent sends automatically to any lead where a public number was found.
- Instagram: approved message is placed in the Instagram Manual Send section — a human sends it manually and marks it as Sent, which enters it into the CRM pipeline.
- Every send logged with: timestamp, account used, platform, message sent.

Task 14 — Hot lead instant WhatsApp alert

- Any lead scoring 8-10 (Hot) triggers an immediate WhatsApp alert to BOTH configured numbers — fired as soon as scoring is complete, not waiting for the end of the run.
- Alert includes: business name, platform, score, written reasoning, and a direct link to the lead's record in the dashboard.

Task 15 — CRM record creation and pipeline entry

- Every lead that is contacted (LinkedIn sent, WhatsApp sent, or Instagram manually marked as sent) becomes a full CRM record.
- Pipeline stages: New → Contacted → Replied → Interested → Meeting Booked → Deal Closed → Lost
- Agent moves leads forward automatically when it detects a reply.
- Hussein and Mahmoud can move leads manually at any time.
- Meeting Booked and Deal Closed are MANUALLY logged by Hussein/Mahmoud — the automated appointment setter is out of scope for this build but the pipeline stages exist so the team can track meetings and deals manually. The dashboard counts them.
- Every stage change is timestamped for a full audit trail.

Task 16 — Follow-up and re-engagement execution

- For leads where Hussein has enabled follow-up and set a time: agent sends the follow-up at exactly that time.
- Follow-up is a PER-LEAD decision — not a blanket schedule for all leads.
- Hussein sets the timing himself per lead — the agent does not decide the delay.
- RE-ENGAGEMENT: if a lead replied once then went cold (replied but stopped responding), the agent can send a re-engagement message. This is treated as a follow-up type and is subject to the same per-lead on/off and timing controls.
- Follow-up and re-engagement pause automatically the moment the lead replies again.
- Second contact (re-contact): the agent can contact the same lead a maximum of TWO times total. Hussein sets the time gap before the second contact (his example: 1 month — but fully his choice, set per lead). Hard maximum of 2 contacts.

Task 17 — Daily run completion and summary

- At the end of each daily run: sends a WhatsApp summary to BOTH recipient numbers.
- Summary includes: total leads found (by platform), total messages sent (by platform), hot/warm/cold breakdown, any account warnings that occurred, exact finish time, total run duration.
- Finish time is also displayed visibly in the dashboard run status panel.

7. Dashboard — Every Screen and Section

The dashboard is a responsive PWA (Progressive Web App) working on laptop and on Mahmoud's iPhone (add to home screen via Safari). Every section below must be built.

7.1 Daily Live Feed

Today's leads

- Every lead found and analyzed today shown as individual cards.
- Each card shows: business name, platform badge (LinkedIn/Instagram/WhatsApp), profile photo, follower/headcount count, industry, score (1-10), temperature badge (Hot/Warm/Cold with color coding), ALL weak points found (full list, not hidden — explicitly required), AI opportunities list, founder flag if detected (with two-part notification: company messaged + founder found), the AI-generated message (copyable with one tap), and current pipeline status.
- Hot leads visually highlighted and pinned to the top of the feed.
- Filter cards by: platform, score range, temperature, date.
- Click any card to open the full Lead Detail page.

7.2 Instagram Manual Send Section

Dedicated IG section — specific requirement from planning

- A clearly separated section (tab or distinct panel) showing only today's Instagram leads that are queued for manual sending — AFTER Mahmoud's approval.
- Each entry shows: business name, score, temperature, the AI-generated message (copyable), founder flag if present, and a 'Mark as Sent' button.
- When friend taps 'Mark as Sent': the lead enters the CRM pipeline as Contacted — identical to how an automated LinkedIn send is logged.
- The section also shows leads pending (not yet sent) vs sent today.
- Accessible and fully usable on mobile (Mahmoud's iPhone, Hussein's phone).

7.3 Mahmoud's Approval Section

Daily message approval — hard requirement

- A dedicated tab/section where Mahmoud sees every message queued for today.
- Each queued item shows: business name, platform, score, temperature, the full generated message text, and Edit + Approve buttons.
- Mahmoud can edit the message text before approving.
- Individual approve per message AND an 'Approve All' button for speed.
- Count badge: e.g. '14 messages awaiting approval' visible at a glance.
- Nothing sends until this section is cleared by Mahmoud.
- Fully mobile-friendly — Mahmoud uses this on his iPhone.

7.4 CRM Pipeline Board

Visual kanban pipeline

- Columns: New / Contacted / Replied / Interested / Meeting Booked / Deal Closed / Lost
- Meeting Booked and Deal Closed are manually moved by Hussein/Mahmoud (not automated in this build) — but they are tracked and counted in the analytics.
- Each card shows: name, platform, score, temperature, last activity date.
- Drag and drop to move leads manually between stages.
- Agent moves leads automatically on reply detection.
- Click any card to open the full Lead Detail page.

7.5 Lead Detail Page

Full record for any lead

- Opens on click from anywhere in the dashboard.
- Shows everything: full profile data, website, all social platforms active on, complete weak points list, AI opportunities list, founder name and flag, WhatsApp number if found, score with full written reasoning, temperature, all messages ever sent to this lead with dates, all replies received with dates, current pipeline stage with full stage history and timestamps, follow-up settings (on/off, next scheduled follow-up), contact count (1st or 2nd touch), and a free-text notes field editable by Hussein or Mahmoud.
- Edit follow-up toggle and timing directly from this page.
- Manually move this lead to any pipeline stage from this page.

7.6 Client History

Permanent database of every analyzed lead — answers the original 'previous data' requirement

- Holds every lead ever analyzed — not just today's. Permanent, never resets.
- This directly fulfills the original requirement: 'a section showing previous data — how many clients did we reach, how many deals, how many meetings, how many hot/warm/cold'
- Searchable and filterable by: name, industry, platform, score, temperature, date range, contacted/not-contacted, founder-detected yes/no, stage.
- Each entry opens to the full Lead Detail page.
- Shows totals at the top: total ever analyzed, total contacted, total hot/warm/cold.

7.7 Analytics & KPI Dashboard

All numbers and metrics — answers 'how many did we reach, book, deal, hot/warm/cold'

- HEADLINE KPI CARDS (top row, always visible):
 - Total businesses reached — all time / this month / this week / today
 - Total messages sent — all time and by period
 - Total replies received + reply rate (%)
 - Total meetings booked (manually logged)
 - Total deals closed (manually logged)
 - Total hot leads found
- BREAKDOWN CHARTS:
 - Reached per platform: LinkedIn vs Instagram vs WhatsApp (bar chart)
 - Reached per account: which of the 3 accounts contributed what (bar chart)
 - Messages sent vs pending approval (status indicator)
 - Leads per pipeline stage: count in each stage (funnel chart)
 - Hot / Warm / Cold distribution (donut chart)
 - Reply rate per platform (bar chart)
 - Founders detected (count) + founders manually contacted (count)
 - Instagram-specific: IG leads found / IG DMs sent manually / IG DMs pending send
 - Daily activity timeline: leads found, messages sent, replies per day (line chart)
- FILTERS on all metrics: date range (today/week/month/all time), platform, account
- All numbers update in real time as the agent runs.

7.8 Account Health Monitor

Per-account status panel

- All 3 accounts shown with current status: Active / Warned / Paused.
- For any warned/paused account: shows specific WARNING TYPE and REASON — not generic.
- Manual toggle per account: 'Redistribute this account's quota?' — Hussein's decision, not automatic.
- Per-account daily usage: leads found and messages sent today vs the configured limit.
- Last active time per account.
- Proxy status per account: connected / missing / error.

7.9 Run Status Panel

Live run tracking

- Status indicator: Running / Idle / Error.
- Shows which account is currently active and what task it is performing.
- Exact start time and finish time of each day's run.
- Total run duration (e.g. 'Completed in 2h 14m').
- Progress indicator: X of 150 leads processed today.

8. Dashboard Settings — Every Editable Control

Every setting below is editable from the dashboard without touching any code. Changes take effect on the next run unless stated otherwise.

Targeting

- Target niche / industry / business type
- Target location / country / market (geography-based)
- Outreach language(s)

Volume

- Instagram leads per day per account (default: 20)
- LinkedIn leads per day per account (default: 30)

Scheduling — per account individually

- Run time for Account 1
- Run time for Account 2
- Run time for Account 3
- (Each is set independently — not one shared time for all three)

Message style

- Active message style: Direct / Discovery / (others can be added)
- Style duration: how many days/weeks the current style stays active before rotating
- (Both are editable anytime — change takes effect from next message generation batch)

Contact and follow-up rules

- Per-lead follow-up: on/off toggle + timing — set from each lead's detail page
- Re-contact gap: minimum time before a second touch is allowed per lead
- Contact cap is fixed at 2 max per lead — not user-editable (hard rule)

Account and proxy configuration

- Proxy IP / credentials for Account 1 — editable slot, can be empty at launch
- Proxy IP / credentials for Account 2 — editable slot, can be empty at launch
- Proxy IP / credentials for Account 3 — editable slot, can be empty at launch
- Account redistribution: when an account is paused, Hussein actively toggles whether its quota goes to the healthy accounts — never automatic

WhatsApp notification settings

- Recipient number 1 (e.g. Hussein)
- Recipient number 2 (e.g. Mahmoud)
- Both numbers receive all alert types by default: hot lead alerts, daily summary, account warnings, founder-found alerts, approval reminders
- Optional: route specific alert types to specific numbers (e.g. approval reminders to Mahmoud only) — routing logic to be finalized during build

9. WhatsApp Notifications — All Alert Types

Every notification the agent sends

- HOT LEAD ALERT — immediate, fires when any lead scores 8-10, to BOTH numbers. Includes: business name, platform, score, reasoning, dashboard link.
- FOUNDER FOUND ALERT — fires when a founder is detected alongside a company message, to BOTH numbers. Includes: business name, founder name, platform, dashboard link.
- ACCOUNT WARNING ALERT — fires immediately when an account shows a problem, to BOTH numbers. Includes: which account, warning type, specific reason.
- APPROVAL REMINDER — fires to Mahmoud's number when the day's messages are queued and awaiting his approval. Includes count of messages waiting.
- DAILY RUN SUMMARY — fires at run completion to BOTH numbers. Includes: total leads found (by platform), total messages sent (by platform), hot/warm/cold count, any warnings that occurred, exact finish time, total run duration.

10. CRM, Pipeline & Follow-Up — Full Specification

Every field in a CRM record

- Business name, platform discovered on, profile URL, profile photo
- Industry, company size estimate, website URL
- All social media platforms the business is active on
- Complete weak points list (every weak point found, not summarized)
- AI opportunities list (specific to their business)
- Founder detected: yes/no + founder name and bio phrase found
- WhatsApp number found: yes/no + number
- Lead score (1-10) + temperature (Hot/Warm/Cold) + full written reasoning
- Generated message text
- Send date, account used, platform sent on
- All replies received with timestamps
- Pipeline stage + complete stage history with timestamps
- Follow-up: on/off + next scheduled follow-up date
- Contact count: 1st touch or 2nd touch
- Date of first contact
- Notes field: free text, editable by Hussein or Mahmoud at any time

Pipeline stages

- New → Contacted → Replied → Interested → Meeting Booked → Deal Closed → Lost
- Agent moves leads automatically on reply detection (New→Contacted on send, Contacted→Replied on reply).
- All other moves are manual — drag on the board or change from lead detail page.
- Meeting Booked and Deal Closed are manually logged — no automated booking in this build. But they ARE tracked and counted in analytics to answer the original requirement of 'how many deals, how many meetings.'
- Every stage move is timestamped.

Follow-up and re-engagement rules

- Follow-up is a per-lead decision — Hussein sets it individually.
- Hussein sets the exact timing per lead — not a preset delay.
- Re-engagement: when a lead replied once then went cold, the agent can send a re-engagement message. Treated as a follow-up type under the same per-lead controls.
- Maximum 2 total contacts per lead (first outreach + one re-contact). Hard limit.
- Re-contact gap: Hussein sets the minimum time before the second contact per lead.
- All follow-up and re-engagement pause instantly when the lead replies.

11. Account Safety & Proxy Plan

The shared-phone origin issue

- All 3 accounts were created from the same phone/laptop — platforms logged this link.
- Fix going forward: give each account independent behavior. Hussein's phone must not log into these accounts casually anymore — all activity goes through the agent's dedicated browser sessions via their individual proxy IPs.
- Over time, consistent independent behavior per account makes the original shared origin irrelevant.

Proxy types — what they are and which to use

- A proxy changes the IP address traffic appears to come from — its only job is to prevent accounts from being flagged or banned for suspicious activity.
- DATACENTER proxies: IPs from cloud servers. Immediately detected by Instagram and LinkedIn. DO NOT USE for this system.
- RESIDENTIAL proxies: IPs from real home internet connections. Look like genuine users. ~\$1-5/GB. The standard safe choice. Use for LinkedIn.
- MOBILE proxies: IPs from real phone carrier networks (4G/5G). Hardest to detect, safest for Instagram specifically. ~\$8+/GB or a SIM plan.
- STICKY proxies (required): the SAME IP for the entire account session (login to logout) — looks like one real person. This is what must be used.
- ROTATING proxies (wrong): different IP on every request — looks like a bot. Never use.

Per-account IP assignment — critical rule

- Each account gets its OWN dedicated sticky IP — never shared:
- Account 1 → Proxy IP #1 (always, permanently assigned)
- Account 2 → Proxy IP #2 (always, permanently assigned)
- Account 3 → Proxy IP #3 (always, permanently assigned)
- The agent's browser sessions must NEVER mix these — each account's context is fully isolated on its own proxy.
- 3 dedicated residential sticky IPs needed total (one per account).

Proxy cost and build plan

- At ~150 profiles/day (~10-15 GB/month total), cost with a budget provider: ~\$5-30/month.
- Recommended pay-as-you-go providers (non-expiring credits): DataImpulse (~\$1/GB), Evomi (~\$0.49/GB), IPRoyal (~\$1.75/GB). Avoid any datacenter-only provider.
- BUILD PLAN: launch with NO proxy during low-volume warm-up (free, safe at 5-10/day). Add proxy credentials in dashboard settings when scaling up. Zero code changes needed — purely a settings update.
- Proxy config slot per account is built into the agent from day one, even if empty.

Warm-up plan — do NOT start at full volume

- Do NOT start at 50/day per account on day one.
- Start at 5-10 actions/day per account. Increase gradually over 2-4 weeks to reach the steady-state 20 Instagram + 30 LinkedIn per account.
- Vary timing and pattern slightly between the 3 accounts — identical parallel behavior at the same times re-creates a 'one operator' signal even with separate IPs.
- Consider building an automatic warm-up schedule: each account starts at a low limit and increments weekly until reaching the configured maximum.

12. Technical Architecture

Platform: responsive PWA

- Must work on BOTH laptop (Hussein) and iPhone (Mahmoud).
- Decision: one responsive web app (PWA) — one codebase, any device, any browser.
- On Mahmoud's iPhone: Safari → Add to Home Screen → works like a native app, full-screen, with icon. No App Store, no Apple developer account needed.
- Alerts via WhatsApp — no dependency on PWA push notification support.

Confirmed tech stack

- Frontend: React + Tailwind CSS (Hussein's existing stack), fully responsive
- Backend / agent: Python
- Agent brain: Claude API
- Haiku model: analysis, scoring, qualification, founder detection (cost-efficient)
- Sonnet model: message generation (better writing quality, drives reply rates)
- Database: Supabase (Postgres + Auth + Realtime) — FREE TIER. Only 2 users (Hussein + Mahmoud), well within free limits.
- Dashboard hosting: Vercel free tier
- Agent server: needs always-on process for scheduled daily runs.
- Option A: Oracle Cloud free tier (genuinely free, always-on)
- Option B: cheap VPS ~\$5/month (e.g. Hetzner) — more reliable for production
- VPS recommended for a real production system
- Proxy: config slot per account built from day one, starts empty

13. Running Cost Estimate

Monthly operating costs

- Claude API (Haiku analysis + Sonnet messages, ~4,500 leads/month): \$35-55/month
- Supabase: \$0 (free tier, 2 users)
- Dashboard hosting (Vercel): \$0
- Agent server: \$0 (Oracle free) or ~\$5/month (VPS)
- Proxies (~10-15 GB/month, budget residential): \$5-30/month
- WhatsApp notifications (internal team, low volume): \$0-15/month
- TOTAL REALISTIC RANGE: \$50-105/month, likely ~\$70-90/month
- Retainer of \$150-200/month from client covers costs with healthy margin

14. Recommended Build Order for Claude Code

Step-by-step build sequence

- Step 1: Supabase schema — leads, CRM/pipeline, client history, settings, accounts, notifications log tables
- Step 2: Agent core — account pool manager with per-account proxy slots, LinkedIn discovery, Instagram discovery, qualification logic
- Step 3: Claude analysis pipeline — deep analysis, founder detection, WhatsApp number detection, lead scoring (Haiku), message generation (Sonnet), personalization
- Step 4: Approval system — message queuing, Mahmoud's approval section, send gate
- Step 5: Sending layer — LinkedIn auto-send, WhatsApp auto-send, Instagram manual-send section with Mark as Sent flow
- Step 6: CRM and pipeline — record creation, automatic stage movement on reply, follow-up scheduler, re-engagement logic, re-contact tracker (max 2)
- Step 7: WhatsApp notifications — all 5 alert types to both numbers
- Step 8: Dashboard frontend — live feed, IG manual section, approval section, pipeline board, lead detail, client history, analytics/KPIs, account health monitor, run status panel, full settings panel
- Step 9: Deploy — dashboard to Vercel, agent to always-on server
- Step 10: Test at low volume with no proxy → add proxies in settings → warm-up ramp

15. Explicitly Out of Scope — Reference Only

The following were explored in detail during planning and then deliberately excluded when the decision was made to build the outreach agent only. Listed so they are not accidentally added. All can be revisited as a future Phase 2.

Advanced tier features — not being built now

- Automated appointment setter (Calendly integration, booking link auto-send)
- Meeting-booked WhatsApp alerts (automated — manual logging exists in pipeline)
- Meeting preparation agent (pre-call brief auto-delivered 1 hour before)
- Call summary agent (transcript/recording → auto CRM update)
- Proposal generator (auto scope/pricing/ROI/timeline document)
- Post-meeting follow-up sequences (Day 1/3/5/7/14 automated nurture)
- Customer support mode (inbound message handling, knowledge base, escalation routing)
- Industry-wide campaign blasting ('message all restaurants with our offer')
- Deep company intelligence beyond Section 6 (Meta Ad Library ad-spend detection)
- Multi-version messaging (3 tone variations generated per lead)
- Smart context-aware re-engagement referencing live recent activity (basic re-engagement under follow-up controls IS in scope — see Task 16)